

# Assignment: Build an App for AI Agents

**Course:** MIT — Building with AI Agents **Due:** See Canvas for deadline **Format:** Individual or teams of 2

---

## Overview

You will build a **web application that AI agents can use autonomously**. Your app will expose a set of API endpoints and a `skill.md` file that teaches any OpenClaw agent how to interact with it — without human intervention.

**You can build any app you want.** The only requirement is that agents can discover it, learn it, and use it through the skill.md protocol. Some ideas:

- **Study group finder** — agents find classmates studying the same topics
- **Event planner** — agents coordinate meetups, RSVPs, and schedules
- **Book/movie club** — agents recommend and discuss media with each other
- **Marketplace** — agents post and browse listings on behalf of their humans
- **Debate forum** — agents argue positions on topics and vote
- **Recipe exchange** — agents share and rate recipes from their humans
- **Fitness challenge tracker** — agents log workouts and compete
- **Research paper matchmaker** — agents find collaborators based on research interests
- **Lost & found board** — agents post and search for lost items
- **Confession wall** — agents post anonymous confessions and react to others

Or anything else. Get creative.

---

## What You're Building

Your app has three parts:

### 1. Protocol Files (how agents discover and learn your app)

These are markdown/JSON files served at specific URLs that teach agents what your app does and how to use it.

File	URL	Purpose
<b>skill.md</b>	<code>/skill.md</code>	Complete API documentation — teaches agents every endpoint, with examples
<b>heartbeat.md</b>	<code>/heartbeat.md</code>	Task loop — tells agents what to do and when they're done
<b>skill.json</b>	<code>/skill.json</code>	Package metadata — name, version, description, emoji

**skill.md** is the most important file. Think of it as a user manual written for AI instead of humans. When an OpenClaw agent reads your `skill.md`, it should be able to start using your app immediately — registering, authenticating, and calling every endpoint without any human help.

**heartbeat.md** is a task loop that drives the agent forward. It's not a passive check-in — it tells the agent: "Here's what you need to accomplish. Keep going until you're done. If something goes wrong, ask your human."

**skill.json** is simple metadata so agent platforms can display your app's name, description, and emoji.

## 2. Backend API (what agents actually call)

A set of REST API endpoints that agents interact with. At minimum, your app needs:

- **Agent registration** — an endpoint where agents register themselves and get an API key
- **Agent claiming** — a way for humans to claim/verify ownership of their agent
- **Core functionality** — whatever your app does (posting, browsing, messaging, voting, etc.)
- **Bearer token auth** — every request (except registration) requires an API key

## 3. Frontend (what humans see)

A web interface where humans can:

- See what's happening in the app (browse content, view activity)
- Claim their agent (click a link, that's it)
- View results/output of whatever your app does

The frontend doesn't need to be fancy, but it should be functional and look decent.

---

## Step-by-Step Guide

### Step 1: Set Up Your Project

**Use whatever tech stack you want.** The example below uses Next.js + MongoDB, but you can use Flask + SQLite, Express + PostgreSQL, Django + Supabase, Rails, Go — whatever you're comfortable with. The only thing that matters is that your app serves `skill.md`, `heartbeat.md`, `skill.json`, has API endpoints, and has a frontend.

**Example with Next.js + MongoDB** (recommended if you don't have a strong preference):

```
npx create-next-app@latest my-agent-app --typescript --app --tailwind  
cd my-agent-app  
npm install mongoose nanoid
```

**Pick any database you want.** Some free options:

Database	Free Tier	Good for
<a href="#">MongoDB Atlas</a>	512MB free forever	Document-based, flexible schemas
<a href="#">Supabase</a>	500MB, 2 projects free	PostgreSQL, built-in auth
<a href="#">PlanetScale</a>	1 DB free	MySQL, great for relational data
<a href="#">Turso</a>	9GB free	SQLite at the edge, simple
<a href="#">Neon</a>	512MB free	Serverless PostgreSQL
SQLite file	Unlimited, free	Simplest possible, works locally

If using MongoDB Atlas:

1. Go to [cloud.mongodb.com](#) → create a free M0 cluster
2. Create a database user → get connection string

3. Replace `<password>` in the URI with your actual password

Create `.env.local` (or equivalent for your framework):

```
MONGODB_URI=mongodb+srv://username:password@cluster.mongodb.net/?  
retryWrites=true&w=majority  
MONGODB_DB=your-app-name  
APP_URL=http://localhost:3000  
NEXT_PUBLIC_APP_URL=http://localhost:3000  
ADMIN_KEY=pick-any-secret-string
```

**Important:** Never push credentials to GitHub. Add `.env*.local` to `.gitignore`.

## Step 2: Database Connection

(Skip this if you're using a different database — adapt to your ORM/driver.)

If using MongoDB with Mongoose, create `lib/db/mongodb.ts` — this handles connection pooling for serverless environments:

```
import mongoose from 'mongoose';

const MONGODB_URI = process.env.MONGODB_URI!;
const MONGODB_DB = process.env.MONGODB_DB || 'my-agent-app';

if (!MONGODB_URI) throw new Error('Missing MONGODB_URI');

let cached = (global as any).mongoose;
if (!cached) cached = (global as any).mongoose = { conn: null, promise: null };

export async function connectDB() {
  if (cached.conn) return cached.conn;
  if (!cached.promise) {
    cached.promise = mongoose.connect(MONGODB_URI, { dbName: MONGODB_DB });
  }
  cached.conn = await cached.promise;
  return cached.conn;
}
```

## Step 3: Define Your Models

(Adapt this to your database. The schema is what matters, not the ORM.)

At minimum you need an **Agent** model. Here's an example with Mongoose — if you're using Prisma, Drizzle, SQLAlchemy, or raw SQL, just create the equivalent table/schema. Create `lib/models/Agent.ts`:

```
import mongoose, { Schema, Document } from 'mongoose';

export interface IAgent extends Document {
  name: string;
  description: string;
```

```

apiKey: string;
claimToken: string;
claimStatus: 'pending_claim' | 'claimed';
ownerEmail?: string;
lastActive: Date;
}

const AgentSchema = new Schema<IAgent>({
  name: { type: String, required: true, unique: true },
  description: { type: String, required: true },
  apiKey: { type: String, required: true, unique: true },
  claimToken: { type: String, required: true, unique: true },
  claimStatus: { type: String, default: 'pending_claim' },
  ownerEmail: String,
  lastActive: { type: Date, default: Date.now },
}, { timestamps: true });

export default mongoose.models.Agent || mongoose.model<IAgent>('Agent',
AgentSchema);

```

Then define models for whatever your app does — posts, events, reviews, messages, etc.

---

#### Step 4: Helper Utilities

Create `lib/utils/api-helpers.ts` with reusable functions:

```

import { NextResponse } from 'next/server';
import { nanoid } from 'nanoid';

// Standard success response
export function successResponse(data: any, status = 200) {
  return NextResponse.json({ success: true, data }, { status });
}

// Standard error response
export function errorResponse(error: string, hint: string, status: number) {
  return NextResponse.json({ success: false, error, hint }, { status });
}

// Generate API key for agents
export function generateApiKey(): string {
  return `yourapp_${nanoid(32)}`;
}

// Generate claim token
export function generateClaimToken(): string {
  return `yourapp_claim_${nanoid(24)}`;
}

// Extract API key from Authorization header
export function extractApiKey(header: string | null): string | null {

```

```
if (!header) return null;
return header.replace('Bearer ', '').trim() || null;
}
```

Install nanoid: npm install nanoid

## Step 5: Build Your API Routes

Registration: app/api/agents/register/route.ts

This is the first endpoint any agent calls. It creates an agent and returns an API key.

```
import { NextRequest } from 'next/server';
import { connectDB } from '@/lib/db/mongodb';
import Agent from '@/lib/models/Agent';
import { successResponse, errorResponse, generateApiKey, generateClaimToken } from
'@/lib/utils/api-helpers';

export async function POST(req: NextRequest) {
  await connectDB();
  const { name, description } = await req.json();

  if (!name || !description) {
    return errorResponse('Missing fields', 'Both "name" and "description" required',
400);
  }

  // Check if name is taken
  const existing = await Agent.findOne({ name: new RegExp(`^${name}$`, 'i') });
  if (existing) {
    return errorResponse('Name taken', 'Choose a different name', 409);
  }

  const apiKey = generateApiKey();
  const claimToken = generateClaimToken();
  const baseUrl = process.env.APP_URL || process.env.NEXT_PUBLIC_APP_URL ||
'http://localhost:3000';

  await Agent.create({ name, description, apiKey, claimToken });

  return successResponse({
    agent: {
      name,
      api_key: apiKey,
      claim_url: `${baseUrl}/claim/${claimToken}`,
    },
    important: 'SAVE YOUR API KEY! You cannot retrieve it later.',
  }, 201);
}
```

Claim page: app/claim/[token]/page.tsx

A simple page where the human clicks to claim their agent. No complicated verification — just click and done.

### Auth middleware pattern

For all other endpoints, extract and validate the API key:

```
const apiKey = extractApiKey(req.headers.get('authorization'));
if (!apiKey) return errorResponse('Missing API key', 'Include Authorization header', 401);

const agent = await Agent.findOne({ apiKey });
if (!agent) return errorResponse('Invalid API key', 'Agent not found', 401);
```

### Your app's endpoints

Build whatever endpoints your app needs. For example, if you're building an event planner:

- POST /api/events — create an event
- GET /api/events — list events
- POST /api/events/:id/rsvp — RSVP to an event
- GET /api/events/:id — get event details

Every endpoint should follow this pattern:

- Authenticate via Bearer token
- Do the thing
- Return { success: true, data: {...} } or { success: false, error: "...", hint: "..." }

---

## Step 6: Write Your skill.md

This is the most important part. Create `app/skill.md/route.ts`:

```
import { NextResponse } from 'next/server';

export async function GET() {
  const baseUrl = process.env.APP_URL || process.env.NEXT_PUBLIC_APP_URL ||
  'http://localhost:3000';

  const markdown = `---
name: your-app-name
version: 1.0.0
description: One sentence describing what your app does.
homepage: ${baseUrl}
metadata: {"openclaw":
{"emoji":"⌚","category":"social","api_base":"${baseUrl}/api"}}
---` 

  # Your App Name

  Describe what your app does in 2-3 sentences.
```

```

## Step 1: Register

\\\\`\\\\`\\\\`bash
curl -X POST ${baseUrl}/api/agents/register \
-H "Content-Type: application/json" \
-d '{"name": "YourAgentName", "description": "What you do"}'
\\\\`\\\\`\\\\`


Response: { "data": { "agent": { "api_key": "...", "claim_url": "..." } } }

Save your api_key. Send the claim_url to your human.

## Step 2: Get Claimed

Your human clicks the claim link. Done.

## Step 3: [Your core functionality]

Document every endpoint with:
- What it does
- The curl command to call it
- What the response looks like

...continue for all your endpoints...

## Authentication

All requests (except register) require:
\\\\`\\\\`\\\\`\\
Authorization: Bearer YOUR_API_KEY
\\\\`\\\\`\\\\`\\

## Response Format

Success: { "success": true, "data": {...} }
Error: { "success": false, "error": "...", "hint": "..." }
`;

return new NextResponse(markdown, {
  headers: { 'Content-Type': 'text/markdown; charset=utf-8' },
});
}

```

#### Tips for writing good skill.md:

- **Be explicit.** Don't assume the agent knows anything. Show exact curl commands.
- **Show response examples.** Agents need to know what to expect back.
- **Use numbered steps.** Walk the agent through the flow from registration to completion.
- **Include error hints.** Tell agents what to do if something goes wrong.
- **Mention escalation.** Tell agents: "If you don't know something about your human, message them and ask."

## Step 7: Write Your heartbeat.md

Create `app/heartbeat.md/route.ts`. The heartbeat is a **continuous task loop** — not a passive check-in. It tells the agent what to accomplish and keeps it going until the task is done.

```
import { NextResponse } from 'next/server';

export async function GET() {
  const baseUrl = process.env.APP_URL || process.env.NEXT_PUBLIC_APP_URL ||
  'http://localhost:3000';

  const markdown = `# Your App – Heartbeat

**Keep running this loop until you've completed the goal.**

If anything goes wrong, message your human and ask what to do.

## Your Goal

[Describe what "done" looks like. For example:]
- Post at least 3 items
- Respond to all incoming messages
- Browse and interact with at least 5 other agents' content

## Loop

### Step 1: Check your setup
[Am I registered? Claimed? Profile created?]

### Step 2: Handle incoming activity
[Check for new messages, requests, notifications]

### Step 3: Do the main thing
[Post content, browse others, interact, etc.]

### Step 4: Check if done
[Have I met the goal? If yes, tell my human. If no, go back to Step 2.]

## Error Handling

If something fails:
1. Message your human and explain what happened
2. Ask them what to do
3. Try again once you have an answer

Never silently fail.
`;

  return new NextResponse(markdown, {
    headers: { 'Content-Type': 'text/markdown; charset=utf-8' },
  })
}
```

```
});  
}
```

---

## Step 8: Write Your skill.json

Create app/skill.json/route.ts :

```
import { NextResponse } from 'next/server';  
  
export async function GET() {  
  const baseUrl = process.env.APP_URL || process.env.NEXT_PUBLIC_APP_URL ||  
  'http://localhost:3000';  
  
  return NextResponse.json({  
    name: 'your-app-name',  
    version: '1.0.0',  
    description: 'One sentence about your app.',  
    homepage: baseUrl,  
    metadata: {  
      openclaw: {  
        emoji: '🌐',  
        category: 'social',  
        api_base: `${baseUrl}/api`,  
      },  
    },  
  });  
}
```

---

## Step 9: Build Your Frontend

Your app needs a web interface. At minimum:

- **Landing page** ( / ) — what the app is, how it works, link to skill.md
- **Claim page** ( /claim/[token] ) — where humans claim their agent
- **Content pages** — whatever your app shows (events, posts, reviews, etc.)

Use Tailwind CSS for styling. Here's a minimal landing page pattern:

```
export default function HomePage() {  
  return (  
    <div className="max-w-4xl mx-auto px-4 py-16 text-center">  
      <h1 className="text-5xl font-bold mb-4">Your App Name</h1>  
      <p className="text-xl text-gray-600 mb-8">  
        What your app does in one sentence.  
      </p>  
      <div className="bg-gray-900 rounded-xl p-6 mb-8">  
        <p className="text-gray-300 mb-2">Tell your OpenClaw agent:</p>  
        <code className="text-green-400 text-lg">  
          Read https://your-url/skill.md  
        </code>  
      </div>  
    </div>
```

```
    </div>
  );
}
```

## Step 10: Environment Variables for Production

**Important:** `NEXT_PUBLIC_*` variables get baked in at build time. For URLs that need to resolve correctly in production, use a non-prefixed variable like `APP_URL` and read it at runtime:

```
// DO THIS – works in production
const baseUrl = process.env.APP_URL || process.env.NEXT_PUBLIC_APP_URL ||
'http://localhost:3000';

// NOT THIS – gets baked as "localhost" at build time
const baseUrl = process.env.NEXT_PUBLIC_APP_URL || 'http://localhost:3000';
```

Add `APP_URL` to your deployment environment set to your production URL.

## Step 11: Deploy to Railway

We'll use [Railway](#) for deployment. (You can also use Vercel, Render, Fly.io, or any other platform — your app just needs to be live at a public URL.)

1. Push your code to GitHub (make sure `.env*.local` is in `.gitignore`)
2. Create an account at [railway.com](#)
3. Click **New Project** → **Deploy from GitHub repo** → select your repo
4. Add environment variables in the Railway dashboard (click on your service → **Variables**):
  - Your database connection string (e.g. `MONGODB_URI`)
  - `APP_URL` — your Railway URL (e.g. `https://my-app.up.railway.app`)
  - `ADMIN_KEY` — a secret string for admin endpoints
  - Any other env vars your app needs
5. Create `railway.json` in your project root:

```
{
  "$schema": "https://railway.com/railway.schema.json",
  "build": { "builder": "NIXPACKS" },
  "deploy": {
    "startCommand": "npm start",
    "restartPolicyType": "ON_FAILURE",
    "restartPolicyMaxRetries": 10
  }
}
```

6. Push to GitHub — Railway auto-deploys on every push

7. Go to your service's **Settings** → **Networking** → **Generate Domain** to get your public URL

#### Verify your deployment:

```
# Check skill.md serves correctly with your production URL
curl https://your-app.up.railway.app/skill.md

# Test registration
curl -X POST https://your-app.up.railway.app/api/agents/register \
-H "Content-Type: application/json" \
-d '{"name": "TestAgent", "description": "Testing deployment"}'
```

Make sure `skill.md` shows your production URL (not `localhost`). If it shows `localhost`, check Step 10 — you probably need to set `APP_URL` in your Railway environment variables.

## Testing Your App

Before submitting, test the full agent flow yourself using curl:

```
# 1. Read skill.md – does it explain everything clearly?
curl https://your-url/skill.md

# 2. Register an agent
curl -X POST https://your-url/api/agents/register \
-H "Content-Type: application/json" \
-d '{"name": "TestAgent", "description": "Test"}'

# 3. Claim the agent (use the claim_url from registration)

# 4. Use your API key to call every endpoint documented in skill.md
# 5. Follow the heartbeat.md loop – can you complete the goal?
```

**Ask yourself:** If I were an AI agent reading `skill.md` for the first time, would I know exactly what to do? If the answer is no, your `skill.md` needs more detail.

## Reference Example: ClawMatchStudio

**This is just one example.** It uses Next.js + MongoDB + Railway, but you can use completely different tools. What matters is the protocol (`skill.md`, `heartbeat.md`, `skill.json`) and that agents can use your app.

- **GitHub:** [github.com/mariagorskikh/homework2\\_example](https://github.com/mariagorskikh/homework2_example)
- **Live app:** [clawmatch.up.railway.app](https://clawmatch.up.railway.app)
- **skill.md:** [clawmatch.up.railway.app/skill.md](https://clawmatch.up.railway.app/skill.md)
- **heartbeat.md:** [clawmatch.up.railway.app/heartbeat.md](https://clawmatch.up.railway.app/heartbeat.md)
- **Assignment doc:** [github.com/mariagorskikh/homework2\\_example/blob/main/assignment.md](https://github.com/mariagorskikh/homework2_example/blob/main/assignment.md)

ClawMatchStudio is a team matching app where agents have conversations with each other to find compatible teammates. Study it for the patterns — the protocol files, the API design, the auth flow, the frontend — but build something different and make it your own.

## Other references:

- **OpenClaw** — the agent framework your agent runs on: [openclaw.com](https://openclaw.com)
  - **Moltbook** — a social network for agents that uses the same skill.md protocol: [moltbook.com](https://moltbook.com)
- 

## FAQ

**Q: Do I have to use Next.js?** A: No. Use whatever you want — Flask, Express, Django, Rails, Go, anything. The example uses Next.js because it bundles frontend + API + deployment nicely, but it's just one option. Pick whatever you're most productive with.

**Q: Do I have to use MongoDB?** A: No. Use any database — PostgreSQL, MySQL, SQLite, Supabase, PlanetScale, Turso, Neon, even a JSON file if your app is simple enough. The example uses MongoDB Atlas because the free tier is generous, but it's just one option.

**Q: Do I have to deploy on Railway?** A: No. Deploy anywhere — Vercel, Render, Fly.io, Netlify, your own VPS, whatever. Your app just needs to be live at a public URL so agents can reach it.

**Q: Do I have to follow the exact code structure from the example?** A: No. The example is one way to do it. What matters is: your app serves `skill.md`, `heartbeat.md`, and `skill.json` at the right URLs, has API endpoints that work, has auth, and has a frontend. How you organize the code internally is up to you.

**Q: Do I need agent-to-agent conversations?** A: No. That's what the example does, but your app can be anything. An event planner where agents RSVP, a review board where agents post ratings, a marketplace — whatever you want. Agents just need to be able to use it through `skill.md`.

**Q: How detailed should skill.md be?** A: Very. Include curl commands for every endpoint, show example responses, explain what to do on errors. The agent has never seen your app before — `skill.md` is the only documentation it gets. If an agent can't figure out how to use your app from `skill.md` alone, it needs more detail.

**Q: What if my agent doesn't know something about my human?** A: That's what OpenClaw channels are for. Your `skill.md` should tell agents: "If you don't know something about your human, message them through your channel (WhatsApp, Telegram, Discord, Slack, OpenClaw chat, etc.) and ask."

**Q: Can I work in a team?** A: Teams of up to 2. Both members should contribute and understand the full codebase.

---

*Good luck, and build something cool. 🦸*